

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv('Churn_Modelling.csv')
df.head()
```

```
Out[2]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	0.00
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86
2	3	15619304	Onio	502	France	Female	42	8	159660.80
3	4	15701354	Boni	699	France	Female	39	1	0.00
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82


```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender               10000 non-null  object
6   Age                  10000 non-null  int64
7   Tenure               10000 non-null  int64
8   Balance              10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard             10000 non-null  int64
11  IsActiveMember       10000 non-null  int64
12  EstimatedSalary       10000 non-null  float64
13  Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [4]: df.describe(include='all')
```

Out[4]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
count	10000.00000	1.000000e+04	10000	10000.000000	10000	10000	10000.000000
unique	NaN	NaN	2932	NaN	3	2	NaN
top	NaN	NaN	Smith	NaN	France	Male	NaN
freq	NaN	NaN	32	NaN	5014	5457	NaN
mean	5000.50000	1.569094e+07	NaN	650.528800	NaN	NaN	38.921800
std	2886.89568	7.193619e+04	NaN	96.653299	NaN	NaN	10.487806
min	1.00000	1.556570e+07	NaN	350.000000	NaN	NaN	18.000000
25%	2500.75000	1.562853e+07	NaN	584.000000	NaN	NaN	32.000000
50%	5000.50000	1.569074e+07	NaN	652.000000	NaN	NaN	37.000000
75%	7500.25000	1.575323e+07	NaN	718.000000	NaN	NaN	44.000000
max	10000.00000	1.581569e+07	NaN	850.000000	NaN	NaN	92.000000



```
In [5]: df.isnull().sum()
```

Out[5]:

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

Dropping Irrelevant Features

```
In [6]: df.columns
```

Out[6]:

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',  
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',  
      'IsActiveMember', 'EstimatedSalary', 'Exited'],  
      dtype='object')
```

```
In [7]: df = df.drop(['RowNumber', 'CustomerId', 'Surname'],axis=1)
```

```
In [8]: df.head()
```

```
Out[8]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

Encoding Categorical Data

```
In [9]: df = pd.get_dummies(df,drop_first = True)
```

```
In [10]: df.head()
```

```
Out[10]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	42	2	0.00	1	1	1	101356
1	608	41	1	83807.86	1	0	1	112588
2	502	42	8	159660.80	3	1	0	113561
3	699	39	1	0.00	2	0	0	93600
4	850	43	2	125510.82	1	1	1	79043


Some insights about the target variable

```
In [11]: df['Exited'].value_counts()
```

```
Out[11]: 0    7963
         1    2037
         Name: Exited, dtype: int64
```

```
In [12]: sns.countplot(df['Exited'])
```

```
Out[12]: <AxesSubplot:xlabel='Exited', ylabel='count'>
```



```
In [13]: X = df.drop('Exited',axis=1)  
y = df['Exited']
```

Handling Imbalanced Data with SMOTE

```
In [14]: from imblearn.over_sampling import SMOTE
```


```
In [15]: X_res, y_res = SMOTE().fit_resample(X,y)
```

```
In [16]: y_res.value_counts()
```

```
Out[16]: 1    7963  
         0    7963  
         Name: Exited, dtype: int64
```

```
In [17]: sns.countplot(y_res)
```

```
Out[17]: <AxesSubplot:xlabel='Exited', ylabel='count'>
```



Splitting The Dataset into Training Set and Test Set

```
In [18]: from sklearn.model_selection import train_test_split
```

```
In [19]: X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.2)
```

Feature Scaling

```
In [20]: from sklearn.preprocessing import StandardScaler
```

```
In [21]: sc = StandardScaler()
```

```
In [22]: X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [23]: X_train
```

```
Out[23]: array([[ 0.75877907, -0.8982394 ,  0.78677282, ..., -0.575693 ,
                  2.16445251, -0.84390722],
                [-0.68402381, -0.7976025 ,  1.15294588, ..., -0.575693 ,
                  -0.4620106 ,  1.18496439],
                [-1.32891298,  0.2087665 ,  0.78677282, ..., -0.575693 ,
                  2.16445251,  1.18496439],
                ...,
                [-0.00634367,  0.0074927 ,  0.42059976, ..., -0.575693 ,
                  -0.4620106 , -0.84390722],
                [ 0.47459063, -0.9988763 , -1.41026556, ..., -0.575693 ,
                  -0.4620106 ,  1.18496439],
                [ 1.63320506,  2.22150449, -1.41026556, ..., -0.575693 ,
                  -0.4620106 , -0.84390722]])
```

Logistic Regression

```
In [24]: from sklearn.linear_model import LogisticRegression
```

```
In [25]: log = LogisticRegression()
```

```
In [26]: log.fit(X_train, y_train)
```

```
Out[26]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [27]: y_pred1 = log.predict(X_test)
```

```
In [28]: from sklearn.metrics import accuracy_score
```

```
In [29]: accuracy_score(y_test, y_pred1)
```

```
Out[29]: 0.7748011720385098
```

```
In [30]: from sklearn.metrics import precision_score, recall_score, f1_score
```

```
In [31]: precision_score(y_test, y_pred1)
```

```
Out[31]: 0.7582957804178615
```

```
In [32]: recall_score(y_test, y_pred1)
```

```
Out[32]: 0.79204107830552
```

```
In [33]: f1_score(y_test,y_pred1)
```

```
Out[33]: 0.7748011720385098
```

SVC

```
In [34]: from sklearn import svm
```

```
In [35]: svm = svm.SVC()
```

```
In [36]: svm.fit(X_train,y_train)
```

```
Out[36]: SVC()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [37]: y_pred2=svm.predict(X_test)
```

```
In [38]: accuracy_score(y_test, y_pred2)
```

```
Out[38]: 0.8287986605274174
```

```
In [39]: precision_score(y_test, y_pred2)
```

```
Out[39]: 0.8214134574693187
```

```
In [40]: recall_score(y_test, y_pred2)
```

```
Out[40]: 0.8305519897304237
```

```
In [41]: f1_score(y_test,y_pred2)
```

```
Out[41]: 0.8259574468085107
```

KNeighbors Classifier

```
In [42]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [43]: knn = KNeighborsClassifier()
```

```
In [44]: knn.fit(X_train,y_train)
```

```
Out[44]: KNeighborsClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [45]: y_pred3=knn.predict(X_test)
```

```
In [46]: accuracy_score(y_test, y_pred3)
```

```
Out[46]: 0.8080786940142319
```

```
In [47]: precision_score(y_test, y_pred3)
```

```
Out[47]: 0.785829307568438
```

```
In [48]: recall_score(y_test, y_pred3)
```

```
Out[48]: 0.8352588789045785
```

```
In [49]: f1_score(y_test,y_pred3)
```

```
Out[49]: 0.8097904998962872
```

Desicion Tree Classifier

```
In [50]: from sklearn.tree import DecisionTreeClassifier
```

```
In [51]: dt = DecisionTreeClassifier()
```

```
In [52]: dt.fit(X_train, y_train)
```

```
Out[52]: DecisionTreeClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [53]: y_pred4=dt.predict(X_test)
```

```
In [54]: accuracy_score(y_test, y_pred4)
```

```
Out[54]: 0.7986605274173294
```



```
In [55]: precision_score(y_test, y_pred4)
```

```
Out[55]: 0.7804977560179519
```

```
In [56]: recall_score(y_test, y_pred4)
```

```
Out[56]: 0.8185708172871202
```

```
In [57]: f1_score(y_test,y_pred4)
```

```
Out[57]: 0.7990810359231411
```

Random Forest Classifier

```
In [58]: from sklearn.ensemble import RandomForestClassifier
```

```
In [59]: rf = RandomForestClassifier()
```

```
In [60]: rf.fit(X_train,y_train)
```

```
Out[60]: RandomForestClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [61]: y_pred5=rf.predict(X_test)
```

```
In [62]: accuracy_score(y_test,y_pred5)
```

```
Out[62]: 0.8503557974047719
```

```
In [63]: precision_score(y_test, y_pred5)
```

```
Out[63]: 0.8407563025210084
```

```
In [64]: recall_score(y_test, y_pred5)
```

```
Out[64]: 0.8562259306803595
```

```
In [65]: f1_score(y_test,y_pred5)
```

```
Out[65]: 0.8484206063175749
```

Gradient Boosting Classifier

```
In [66]: from sklearn.ensemble import GradientBoostingClassifier
```

```
In [67]: gbc = GradientBoostingClassifier()
```

```
In [68]: gbc.fit(X_train, y_train)
```

```
Out[68]: GradientBoostingClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [69]: y_pred6=gbc.predict(X_test)
```

```
In [70]: accuracy_score(y_test,y_pred6)
```

```
Out[70]: 0.8340309753034743
```

```
In [71]: precision_score(y_test,y_pred6)
```

```
Out[71]: 0.8307626392459297
```

```
In [72]: recall_score(y_test,y_pred6)
```

```
Out[72]: 0.8296961916987591
```

```
In [73]: f1_score(y_test,y_pred6)
```

```
Out[73]: 0.8302290730036395
```

XGBoost

```
In [86]: import xgboost as xgb
```

```
model_xgb = xgb.XGBClassifier(random_state=42, verbosity = 0)
model_xgb.fit(X_train, y_train)
```

```
Out[86]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
                      gamma=0, gpu_id=-1, importance_type=None,
                      interaction_constraints='', learning_rate=0.300000012,
                      max_delta_step=0, max_depth=6, min_child_weight=1, missing=nan,
                      monotone_constraints='()', n_estimators=100, n_jobs=8,
                      num_parallel_tree=1, predictor='auto', random_state=42,
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                      tree_method='exact', validate_parameters=1, verbosity=0)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [87]: y_pred7=model_xgb.predict(X_test)
```

```
In [88]: accuracy_score(y_test,y_pred7)
```

```
Out[88]: 0.8522394307241523
```

```
In [89]: precision_score(y_test,y_pred7)
```

```
Out[89]: 0.8382413936126089
```

```
In [90]: recall_score(y_test,y_pred6)
```

```
Out[90]: 0.8296961916987591
```

```
In [91]: f1_score(y_test,y_pred6)
```

```
Out[91]: 0.8302290730036395
```

Accuracy Summary

```
In [92]: performance_summary = pd.DataFrame({
    'Model': ['LR', 'SVC', 'KNN', 'DT', 'RF', 'GBC', 'XGB'],
    'ACC': [accuracy_score(y_test, y_pred1),
            accuracy_score(y_test, y_pred2),
            accuracy_score(y_test, y_pred3),
            accuracy_score(y_test, y_pred4),
            accuracy_score(y_test, y_pred5),
            accuracy_score(y_test, y_pred6),
            accuracy_score(y_test, y_pred7)
    ]
})
```


```
In [93]: performance_summary
```

```
Out[93]:
```

	Model	ACC
0	LR	0.774801
1	SVC	0.828799
2	KNN	0.808079
3	DT	0.798661
4	RF	0.850356
5	GBC	0.834031
6	XGB	0.852239

```
In [94]: sns.barplot(performance_summary['Model'], performance_summary['ACC'])
```

```
Out[94]: <AxesSubplot:xlabel='Model', ylabel='ACC'>
```



As we can see, XGBoost Classifier has highest accuracy

```
In [95]: performance_summary = pd.DataFrame({
    'Model': ['LR', 'SVC', 'KNN', 'DT', 'RF', 'GBC', 'XGB'],
    'PRECISION': [precision_score(y_test, y_pred1),
                  precision_score(y_test, y_pred2),
                  precision_score(y_test, y_pred3),
                  precision_score(y_test, y_pred4),
                  precision_score(y_test, y_pred5),
                  precision_score(y_test, y_pred6),
                  precision_score(y_test, y_pred6)]
    })
```


```
In [96]: performance_summary
```

```
Out[96]:
```

	Model	PRECISION
0	LR	0.758296
1	SVC	0.821413
2	KNN	0.785829
3	DT	0.780498
4	RF	0.840756
5	GBC	0.830763
6	XGB	0.830763

```
In [97]: sns.barplot(performance_summary['Model'], performance_summary['PRECISION'])
```

```
Out[97]: <AxesSubplot:xlabel='Model', ylabel='PRECISION'>
```



Saving the best model, XGBoost

```
In [80]: X_train = sc.fit_transform(X_train)
```

```
In [98]: model_xgb.fit(X_res,y_res)
```

```
Out[98]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
                      gamma=0, gpu_id=-1, importance_type=None,
                      interaction_constraints='', learning_rate=0.300000012,
                      max_delta_step=0, max_depth=6, min_child_weight=1, missing=nan,
                      monotone_constraints='()', n_estimators=100, n_jobs=8,
                      num_parallel_tree=1, predictor='auto', random_state=42,
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                      tree_method='exact', validate_parameters=1, verbosity=0)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [99]: import joblib
```

```
In [100]: joblib.dump(model_xgb, 'churn_predict_model')
```

```
Out[100]: ['churn_predict_model']
```

```
In [101]: model = joblib.load('churn_predict_model')
```

```
In [85]: df.columns
```

```
Out[85]: Index(['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
                'IsActiveMember', 'EstimatedSalary', 'Exited', 'Geography_Germany',
                'Geography_Spain', 'Gender_Male'],
               dtype='object')
```